

Parallel System Architectures — Lab 2

Andrea Di Dio — 12967424

January 26, 2020

1 Design Choices

For this assignment we had to implement a split-transaction bus. My implementation is inspired by the *simple.bus* example from the SystemC documentation. There are two main components which have been added on top of the code for assignment 1 in order to implement the valid-invalid cache coherency protocol. The Bus is non-blocking and works with an arbiter to decide which request to satisfy on the falling edge on the clock. I decided to use a two-phase synchronization approach where the caches (bus masters) make the requests to the bus on the rising edge of the clock and then one request from the queue is satisfied on the falling edge of the clock.

The caches are connected to the bus by means of a port to allow for a TLM implementation of the components such that the caches can use two method calls (*port.bus*→*write()* or *port.bus*→*read()*). These calls place a new *bus_sig_t* request in the queue which represents the type of bus operation, the address and the id of the requester cache. Each cache has a separate *snoop* thread, which waits for the signal driven by the bus which uses it to “reply” to requests to change, and take different actions such as invalidating a cache line whenever a network read is observed or updating the line requested from memory.

The arbiter is the component that at every falling edge of the clock picks a request to be processed. Unless the queue contains a request which is “in process” already and should therefore be prioritized to finish, the arbiter simply picks the oldest request which allows to serve the requests in FIFO order.

The *execute* thread of the bus, tries to prioritize the replies from memory by writing to the signal as soon as the memory component (bus slave) returns the *MEM_DONE* return value. This allows the requester of that memory operation and other snooping caches to observe a change in value as soon as possible.

2 Results — Raw Data

2.1 Debug Tracefile

2.1.1 2 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	30	0	30	25	1	24	1.818182
1	32	0	32	35	1	34	1.492537

Total Memory Accesses = 124

Average Bus Acquisition Time = 2331152 ps

Total Simulation Time = 12625 ns

2.1.2 4 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	8	0	8	8	1	7	6.250000
1	27	0	27	32	0	32	0
2	43	1	42	38	2	36	3.703704
3	45	0	45	42	0	42	0

Total Memory Accesses = 137

Average Bus Acquisition Time = 4747518 ps

Total Simulation Time = 13938 ns

2.1.3 8 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	6	0	6	4	0	4	0
1	34	0	34	22	0	22	0
2	35	0	35	43	0	43	0
3	39	2	37	46	2	44	4.705882
4	36	0	36	55	0	55	0
5	52	0	52	47	0	47	0
6	48	3	45	51	2	49	5.050505
7	42	1	41	55	5	50	6.185567

Total Memory Accesses = 153

Average Bus Acquisition Time = 6667932 ps

Total Simulation Time = 15554 ns

2.2 Random Tracefile

2.2.1 2 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	16496	385	16111	16402	393	16009	2.364885
1	24556	865	23691	24804	877	23927	3.529173

Total Memory Accesses = 85616

Average Bus Acquisition Time = 1610251461 ps

Total Simulation Time = 8647318 ns

2.2.2 4 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	8039	85	7954	8251	89	8162	1.068140
1	20450	0	20450	20221	0	20221	0
2	26553	982	25571	26498	1056	25442	3.841586
3	29600	0	29600	29695	3	29692	0.005059

Total Memory Accesses = 94628

Average Bus Acquisition Time = 3321838086 ps

Total Simulation Time = 9557530 ns

2.2.3 8 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	4113	22	4091	4030	35	3995	0.699988
1	18318	0	18318	18440	0	18440	0
2	25834	21	25813	25470	30	25440	0.099407
3	29549	1257	28274	28801	1180	27621	4.207369
4	31366	2	31364	30548	1	30547	0.004845
5	32015	12	32003	31782	26	31756	0.059564
6	32327	1626	30701	32285	1653	30632	5.074909
7	32536	1687	30849	32584	1771	30813	5.310197

Total Memory Accesses = 97930

Average Bus Acquisition Time = 4244396129 ps

Total Simulation Time = 9891031 ns

2.3 FFT Tracefile

2.3.1 2 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	29313	3423	25890	15417	2013	13404	12.152918
1	29262	3176	26086	14801	1897	12904	11.513061

Total Memory Accesses = 68809

Average Bus Acquisition Time = 1083535306 ps

Total Simulation Time = 6949810 ns

2.3.2 4 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	11274	1749	9525	6553	945	5608	15.111909
1	9417	1462	7955	6083	755	5328	14.303226
2	9834	1524	8310	5523	610	4913	13.895943
3	9881	1507	8374	5972	704	5268	13.946887

Total Memory Accesses = 28751

Average Bus Acquisition Time = 932925367 ps

Total Simulation Time = 2903952 ns

2.3.3 8 Processors

CPU	Reads	RHit	RMiss	Writes	WHit	WMiss	Hirate
0	4956	1103	3853	3154	742	2412	22.749692
1	3983	822	3161	2890	544	2346	19.874873
2	3997	865	3132	2703	502	2201	20.402985
3	4043	858	3185	2695	500	2195	20.154348
4	4055	851	3204	2662	475	2187	19.740956
5	4055	846	3209	2733	510	2223	19.976429
6	4074	826	3248	2710	483	2227	19.295401
7	4124	831	3293	2705	477	2228	19.153610

Total Memory Accesses = 12788

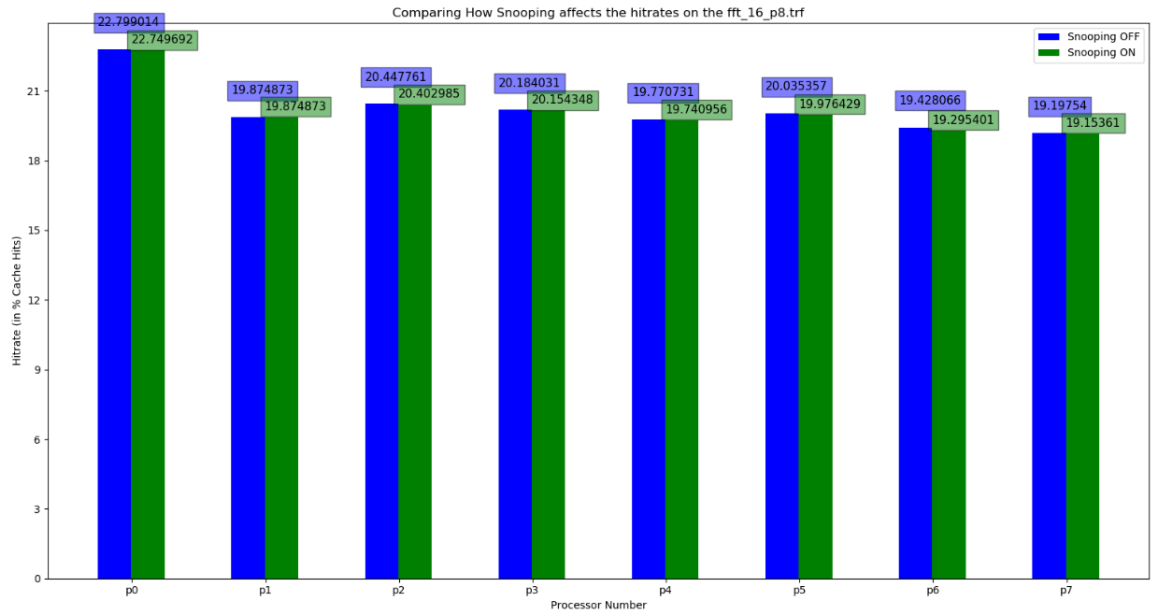
Average Bus Acquisition Time = 533288298 ps

Total Simulation Time = 1291690 ns

3 Experiments

In order to show how the hitrates change by disabling snooping in the caches, I have decided to run the tests only for the 8 processor FFT tracefile as it has the more realistic access patterns.

My hypothesis is that disabling snooping will result in a higher hitrate, albeit being a “fake” hitrate due to the fact that the caches will lose coherency.



The figure above shows that disabling snooping, unexpectedly, does not increase the hitrates by much, but only marginally. One explanation for this could be that the tracefile generates access patterns which are processor independent. I.e., lines that are used by one processor are not used much in other processors, meaning that the writes don’t necessarily affect the validity of the cahce lines in other CPUs.